

dlnd_image_classification

March 28, 2017

1 Image Classification

In this project, you'll classify images from the [CIFAR-10 dataset](#). The dataset consists of airplanes, dogs, cats, and other objects. You'll preprocess the images, then train a convolutional neural network on all the samples. The images need to be normalized and the labels need to be one-hot encoded. You'll get to apply what you learned and build a convolutional, max pooling, dropout, and fully connected layers. At the end, you'll get to see your neural network's predictions on the sample images. ## Get the Data Run the following cell to download the [CIFAR-10 dataset for python](#).

```
In [1]: """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        from urllib.request import urlretrieve
        from os.path import isfile, isdir
        from tqdm import tqdm
        import problem_unittests as tests
        import tarfile

        cifar10_dataset_folder_path = 'cifar-10-batches-py'

        class DLProgress(tqdm):
            last_block = 0

            def hook(self, block_num=1, block_size=1, total_size=None):
                self.total = total_size
                self.update((block_num - self.last_block) * block_size)
                self.last_block = block_num

        if not isfile('cifar-10-python.tar.gz'):
            with DLProgress(unit='B', unit_scale=True, miniters=1, desc='CIFAR-10 Dataset') as p:
                urlretrieve(
                    'https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz',
                    'cifar-10-python.tar.gz',
                    pbar.hook)

        if not isdir(cifar10_dataset_folder_path):
```

```
with tarfile.open('cifar-10-python.tar.gz') as tar:
    tar.extractall()
    tar.close()
```

```
tests.test_folder_path(cifar10_dataset_folder_path)
```

CIFAR-10 Dataset: 171MB [00:52, 3.89MB/s]

All files found!

1.1 Explore the Data

The dataset is broken into batches to prevent your machine from running out of memory. The CIFAR-10 dataset consists of 5 batches, named `data_batch_1`, `data_batch_2`, etc.. Each batch contains the labels and images that are one of the following: * airplane * automobile * bird * cat * deer * dog * frog * horse * ship * truck

Understanding a dataset is part of making predictions on the data. Play around with the code cell below by changing the `batch_id` and `sample_id`. The `batch_id` is the id for a batch (1-5). The `sample_id` is the id for a image and label pair in the batch.

Ask yourself "What are all possible labels?", "What is the range of values for the image data?", "Are the labels in order or random?". Answers to questions like these will help you preprocess the data and end up with better predictions.

```
In [5]: %matplotlib inline
        %config InlineBackend.figure_format = 'retina'

import helper
import numpy as np

# Explore the dataset
batch_id = 3
sample_id = 1000
helper.display_stats(cifar10_dataset_folder_path, batch_id, sample_id)
```

Stats of batch 3:

Samples: 10000

Label Counts: {0: 994, 1: 1042, 2: 965, 3: 997, 4: 990, 5: 1029, 6: 978, 7: 1015, 8: 961, 9: 1029}

First 20 Labels: [8, 5, 0, 6, 9, 2, 8, 3, 6, 2, 7, 4, 6, 9, 0, 0, 7, 3, 7, 2]

Example of Image 1000:

Image - Min Value: 37 Max Value: 230

Image - Shape: (32, 32, 3)

Label - Label Id: 0 Name: airplane



1.2 Implement Preprocess Functions

1.2.1 Normalize

In the cell below, implement the `normalize` function to take in image data, `x`, and return it as a normalized Numpy array. The values should be in the range of 0 to 1, inclusive. The return object should be the same shape as `x`.

```
In [10]: def normalize(x):
         """
         Normalize a list of sample image data in the range of 0 to 1
         : x: List of image data. The image shape is (32, 32, 3)
         : return: Numpy array of normalize data
         """
         # TODO: Implement Function
         return x/255.0

         """
         DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
         """
         tests.test_normalize(normalize)
```

Tests Passed

1.2.2 One-hot encode

Just like the previous code cell, you'll be implementing a function for preprocessing. This time, you'll implement the `one_hot_encode` function. The input, `x`, are a list of labels. Implement the function to return the list of labels as One-Hot encoded Numpy array. The possible values for labels are 0 to 9. The one-hot encoding function should return the same encoding for each value between each call to `one_hot_encode`. Make sure to save the map of encodings outside the function.

Hint: Don't reinvent the wheel.

```
In [29]: def one_hot_encode(x):
        """
        One hot encode a list of sample labels. Return a one-hot encoded vector for each label
        : x: List of sample Labels
        : return: Numpy array of one-hot encoded labels
        """
        # TODO: Implement Function
        # one_hot = np.zeros((len(x), 10))
        # for i, j in enumerate(x):
        #     one_hot[i][j]=1
        one_hot = np.eye(10)[x]
        return one_hot

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        tests.test_one_hot_encode(one_hot_encode)
```

Tests Passed

1.2.3 Randomize Data

As you saw from exploring the data above, the order of the samples are randomized. It doesn't hurt to randomize it again, but you don't need to for this dataset.

1.3 Preprocess all the data and save it

Running the code cell below will preprocess all the CIFAR-10 data and save it to file. The code below also uses 10% of the training data for validation.

```
In [30]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """
        # Preprocess Training, Validation, and Testing Data
        helper.preprocess_and_save_data(cifar10_dataset_folder_path, normalize, one_hot_encode)
```

2 Check Point

This is your first checkpoint. If you ever decide to come back to this notebook or have to restart the notebook, you can start from here. The preprocessed data has been saved to disk.

```
In [102]: """
          DON'T MODIFY ANYTHING IN THIS CELL
          """

          import pickle
          import problem_unittests as tests
          import helper
          import functools
          # Load the Preprocessed Validation data
          valid_features, valid_labels = pickle.load(open('preprocess_validation.p', mode='rb'))
```

2.1 Build the network

For the neural network, you'll build each layer into a function. Most of the code you've seen has been outside of functions. To test your code more thoroughly, we require that you put each layer in a function. This allows us to give you better feedback and test for simple mistakes using our unittests before you submit your project.

Note: If you're finding it hard to dedicate enough time for this course each week, we've provided a small shortcut to this part of the project. In the next couple of problems, you'll have the option to use classes from the [TensorFlow Layers](#) or [TensorFlow Layers \(contrib\)](#) packages to build each layer, except the layers you build in the "Convolutional and Max Pooling Layer" section. TF Layers is similar to Keras's and TFLearn's abstraction to layers, so it's easy to pickup.

However, if you would like to get the most out of this course, try to solve all the problems *without* using anything from the TF Layers packages. You **can** still use classes from other packages that happen to have the same name as ones you find in TF Layers! For example, instead of using the TF Layers version of the conv2d class, [tf.layers.conv2d](#), you would want to use the TF Neural Network version of conv2d, [tf.nn.conv2d](#).

Let's begin!

2.1.1 Input

The neural network needs to read the image data, one-hot encoded labels, and dropout keep probability. Implement the following functions * Implement `neural_net_image_input` * Return a [TF Placeholder](#) * Set the shape using `image_shape` with batch size set to None. * Name the TensorFlow placeholder "x" using the TensorFlow name parameter in the [TF Placeholder](#). * Implement `neural_net_label_input` * Return a [TF Placeholder](#) * Set the shape using `n_classes` with batch size set to None. * Name the TensorFlow placeholder "y" using the TensorFlow name parameter in the [TF Placeholder](#). * Implement `neural_net_keep_prob_input` * Return a [TF Placeholder](#) for dropout keep probability. * Name the TensorFlow placeholder "keep_prob" using the TensorFlow name parameter in the [TF Placeholder](#).

These names will be used at the end of the project to load your saved model.
Note: None for shapes in TensorFlow allow for a dynamic size.

```
In [40]: import tensorflow as tf

def neural_net_image_input(image_shape):
    """
    Return a Tensor for a batch of image input
    : image_shape: Shape of the images
    : return: Tensor for image input.
    """
    # TODO: Implement Function
    return tf.placeholder(tf.float32, (None,)+image_shape, "x")

def neural_net_label_input(n_classes):
    """
    Return a Tensor for a batch of label input
    : n_classes: Number of classes
    : return: Tensor for label input.
    """
    # TODO: Implement Function
    return tf.placeholder(tf.float32, (None, n_classes), "y")

def neural_net_keep_prob_input():
    """
    Return a Tensor for keep probability
    : return: Tensor for keep probability.
    """
    # TODO: Implement Function
    return tf.placeholder(tf.float32, name="keep_prob")

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

tf.reset_default_graph()
tests.test_nn_image_inputs(neural_net_image_input)
tests.test_nn_label_inputs(neural_net_label_input)
tests.test_nn_keep_prob_inputs(neural_net_keep_prob_input)
```

Image Input Tests Passed.
Label Input Tests Passed.
Keep Prob Tests Passed.

2.1.2 Convolution and Max Pooling Layer

Convolution layers have a lot of success with images. For this code cell, you should implement the function `conv2d_maxpool` to apply convolution then max pooling: * Create the weight and bias using `conv_ksize`, `conv_num_outputs` and the shape of `x_tensor`. * Apply a convolution to `x_tensor` using weight and `conv_strides`. * We recommend you use same padding, but you're welcome to use any padding. * Add bias * Add a nonlinear activation to the convolution. * Apply Max Pooling using `pool_ksize` and `pool_strides`. * We recommend you use same padding, but you're welcome to use any padding.

Note: You can't use [TensorFlow Layers](#) or [TensorFlow Layers \(contrib\)](#) for **this** layer, but you can still use TensorFlow's [Neural Network](#) package. You may still use the shortcut option for all the **other** layers.

```
In [180]: def conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides):
    """
    Apply convolution then max pooling to x_tensor
    :param x_tensor: TensorFlow Tensor
    :param conv_num_outputs: Number of outputs for the convolutional layer
    :param conv_ksize: kernal size 2-D Tuple for the convolutional layer
    :param conv_strides: Stride 2-D Tuple for convolution
    :param pool_ksize: kernal size 2-D Tuple for pool
    :param pool_strides: Stride 2-D Tuple for pool
    : return: A tensor that represents convolution and max pooling of x_tensor
    """
    # TODO: Implement Function
    weights_shape = conv_ksize + (x_tensor.shape[3].value, conv_num_outputs)
    weights = tf.Variable(tf.truncated_normal(weights_shape, stddev=0.1), name="conv_weights")
    bias = tf.Variable(tf.zeros(conv_num_outputs), name="conv_bias")

    conv_strides = (1,) + conv_strides + (1,)
    pool_strides = (1,) + pool_strides + (1,)
    pool_size = (1,) + pool_ksize + (1,)

    conv_layer = tf.nn.conv2d(x_tensor, weights, conv_strides, "SAME")
    conv_layer = tf.nn.bias_add(conv_layer, bias)
    conv_layer = tf.nn.relu(conv_layer)
    conv_layer = tf.nn.max_pool(conv_layer, pool_size, pool_strides, "SAME", name="conv_max_pool")
    return conv_layer

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """

tests.test_con_pool(conv2d_maxpool)
```

Tests Passed

2.1.3 Flatten Layer

Implement the `flatten` function to change the dimension of `x_tensor` from a 4-D tensor to a 2-D tensor. The output should be the shape (*Batch Size, Flattened Image Size*). Shortcut option: you can use classes from the [TensorFlow Layers](#) or [TensorFlow Layers \(contrib\)](#) packages for this layer. For more of a challenge, only use other TensorFlow packages.

```
In [181]: def flatten(x_tensor):
    """
    Flatten x_tensor to (Batch Size, Flattened Image Size)
    : x_tensor: A tensor of size (Batch Size, ...), where ... are the image dimensions
    : return: A tensor of size (Batch Size, Flattened Image Size).
    """
    # TODO: Implement Function

    x = tf.convert_to_tensor(x_tensor)
    batch_dim = tf.slice(tf.shape(x), [0], [1])
    x_dim = tf.slice(tf.shape(x), [1], [x.shape.ndims-1])

    flatten_dim = tf.expand_dims(tf.reduce_prod(x_dim), 0)
    flatten_shape = tf.concat([batch_dim, flatten_dim], 0)

    flatten = tf.reshape(x, flatten_shape, name="flatten_layer")

    x_shape = x.shape.as_list()
    batch_dim, spatial_dims = x_shape[0], x_shape[1:]
    if all(spatial_dims):
        flatten.set_shape([batch_dim,
                           functools.reduce(lambda x, y: x * y, spatial_dims)])
    else:
        flatten.set_shape([batch_dim, None])

    # flatten = tf.contrib.layers.flatten(x_tensor)
    return flatten

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """

    tests.test_flatten(flatten)
```

Tests Passed

2.1.4 Fully-Connected Layer

Implement the `fully_conn` function to apply a fully connected layer to `x_tensor` with the shape (*Batch Size, num_outputs*). Shortcut option: you can use classes from the [TensorFlow Layers](#) or

TensorFlow Layers (contrib) packages for this layer. For more of a challenge, only use other TensorFlow packages.

```
In [182]: def fully_conn(x_tensor, num_outputs):
    """
    Apply a fully connected layer to x_tensor using weight and bias
    : x_tensor: A 2-D tensor where the first dimension is batch size.
    : num_outputs: The number of output that the new tensor should be.
    : return: A 2-D tensor where the second dimension is num_outputs.
    """
    # TODO: Implement Function
    weights_shape = (x_tensor.shape[1].value, num_outputs)
    weights = tf.Variable(tf.truncated_normal(weights_shape, stddev=0.1), name="fully_
    bias = tf.Variable(tf.zeros(num_outputs), name="fully_conn_bias")

    fully_conn_layer = tf.matmul(x_tensor, weights)
    fully_conn_layer = tf.nn.bias_add(fully_conn_layer, bias, name="fully_conn_layer")
    fully_conn_layer = tf.nn.relu(fully_conn_layer, name="fully_conn_layer")
    return fully_conn_layer

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """
    tests.test_fully_conn(fully_conn)
```

Tests Passed

2.1.5 Output Layer

Implement the output function to apply a fully connected layer to `x_tensor` with the shape (*Batch Size*, *num_outputs*). Shortcut option: you can use classes from the [TensorFlow Layers](#) or [TensorFlow Layers \(contrib\)](#) packages for this layer. For more of a challenge, only use other TensorFlow packages.

Note: Activation, softmax, or cross entropy should **not** be applied to this.

```
In [185]: def output(x_tensor, num_outputs):
    """
    Apply a output layer to x_tensor using weight and bias
    : x_tensor: A 2-D tensor where the first dimension is batch size.
    : num_outputs: The number of output that the new tensor should be.
    : return: A 2-D tensor where the second dimension is num_outputs.
    """
    # TODO: Implement Function
    weights_shape = (x_tensor.shape[1].value, num_outputs)
    weights = tf.Variable(tf.truncated_normal(weights_shape, stddev=0.1), name="output
    bias = tf.Variable(tf.zeros(num_outputs), name="output_bias")

    output_layer = tf.matmul(x_tensor, weights)
```

```

        output_layer = tf.nn.bias_add(output_layer, bias, name="output_layer")
    return output_layer

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

tests.test_output(output)

```

Tests Passed

2.1.6 Create Convolutional Model

Implement the function `conv_net` to create a convolutional neural network model. The function takes in a batch of images, `x`, and outputs logits. Use the layers you created above to create this model:

- Apply 1, 2, or 3 Convolution and Max Pool layers
- Apply a Flatten Layer
- Apply 1, 2, or 3 Fully Connected Layers
- Apply an Output Layer
- Return the output
- Apply [TensorFlow's Dropout](#) to one or more layers in the model using `keep_prob`.

```

In [186]: def conv_net(x, keep_prob):
    """
    Create a convolutional neural network model
    : x: Placeholder tensor that holds image data.
    : keep_prob: Placeholder tensor that hold dropout keep probability.
    : return: Tensor that represents logits
    """

    # TODO: Apply 1, 2, or 3 Convolution and Max Pool layers
    # Play around with different number of outputs, kernel size and stride
    # Function Definition from Above:
    # conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize)
    conv_num_outputs_list = [16, 32, 64]
    conv_ksize = (4, 4)
    conv_strides = (1, 1)
    pool_ksize = (2, 2)
    pool_strides = (2, 2)

    conv_layer_1 = conv2d_maxpool(x, conv_num_outputs_list[0], conv_ksize, conv_strides, pool_ksize, pool_strides)
    conv_layer_2 = conv2d_maxpool(conv_layer_1, conv_num_outputs_list[1], conv_ksize, conv_strides, pool_ksize, pool_strides)
    conv_layer_3 = conv2d_maxpool(conv_layer_2, conv_num_outputs_list[2], conv_ksize, conv_strides, pool_ksize, pool_strides)

    # TODO: Apply a Flatten Layer
    # Function Definition from Above:
    # flatten(x_tensor)

```

```

flatten_layer = flatten(conv_layer_3)

# TODO: Apply 1, 2, or 3 Fully Connected Layers
#   Play around with different number of outputs
# Function Definition from Above:
#   fully_conn(x_tensor, num_outputs)
fully_conn_layer = fully_conn(flatten_layer, 10)
fully_conn_layer = tf.nn.dropout(fully_conn_layer, keep_prob)
# TODO: Apply an Output Layer
#   Set this to the number of classes
# Function Definition from Above:
#   output(x_tensor, num_outputs)
output_layer = output(fully_conn_layer, 10)

# TODO: return output
return output_layer

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

#####
## Build the Neural Network ##
#####

# Remove previous weights, bias, inputs, etc..
tf.reset_default_graph()

# Inputs
x = neural_net_image_input((32, 32, 3))
y = neural_net_label_input(10)
keep_prob = neural_net_keep_prob_input()

# Model
logits = conv_net(x, keep_prob)

# Name logits Tensor, so that is can be loaded from disk after training
logits = tf.identity(logits, name='logits')

# Loss and Optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))
optimizer = tf.train.AdamOptimizer().minimize(cost)

# Accuracy
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')

```

```
tests.test_conv_net(conv_net)
```

Neural Network Built!

2.2 Train the Neural Network

2.2.1 Single Optimization

Implement the function `train_neural_network` to do a single optimization. The optimization should use `optimizer` to optimize in `session` with a `feed_dict` of the following: * `x` for image input * `y` for labels * `keep_prob` for keep probability for dropout

This function will be called for each batch, so `tf.global_variables_initializer()` has already been called.

Note: Nothing needs to be returned. This function is only optimizing the neural network.

```
In [187]: def train_neural_network(session, optimizer, keep_probability, feature_batch, label_batch):
          """
          Optimize the session on a batch of images and labels
          : session: Current TensorFlow session
          : optimizer: TensorFlow optimizer function
          : keep_probability: keep probability
          : feature_batch: Batch of Numpy image data
          : label_batch: Batch of Numpy label data
          """
          # TODO: Implement Function
          session.run(optimizer, feed_dict={x: feature_batch, y: label_batch, keep_prob: keep_probability})

          """
          DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
          """
          tests.test_train_nn(train_neural_network)
```

Tests Passed

2.2.2 Show Stats

Implement the function `print_stats` to print loss and validation accuracy. Use the global variables `valid_features` and `valid_labels` to calculate validation accuracy. Use a keep probability of 1.0 to calculate the loss and validation accuracy.

```
In [188]: def print_stats(session, feature_batch, label_batch, cost, accuracy):
          """
          Print information about loss and validation accuracy
          : session: Current TensorFlow session
          : feature_batch: Batch of Numpy image data
          : label_batch: Batch of Numpy label data
          """
```

```

: cost: TensorFlow cost function
: accuracy: TensorFlow accuracy function
"""
# TODO: Implement Function
loss = session.run(cost, feed_dict={x: feature_batch, y:label_batch, keep_prob: 1})
valid_accuracy = session.run(accuracy, feed_dict={x: valid_features, y:valid_labels})
print("Loss: {:>6.4f}  Validation Accuracy: {:>4f}".format(loss, valid_accuracy))

```

2.2.3 Hyperparameters

Tune the following parameters: * Set epochs to the number of iterations until the network stops learning or start overfitting * Set batch_size to the highest number that your machine has memory for. Most people set them to common sizes of memory: * 64 * 128 * 256 * ... * Set keep_probability to the probability of keeping a node using dropout

```

In [191]: # TODO: Tune Parameters
epochs = 100
batch_size = 256
keep_probability = 0.8

```

2.2.4 Train on a Single CIFAR-10 Batch

Instead of training the neural network on all the CIFAR-10 batches of data, let's use a single batch. This should save time while you iterate on the model to get a better accuracy. Once the final validation accuracy is 50% or greater, run the model on all the data in the next section.

```

In [170]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

print('Checking the Training on a Single Batch...')
with tf.Session() as sess:
    # Initializing the variables
    sess.run(tf.global_variables_initializer())

    # Training cycle
    for epoch in range(epochs):
        batch_i = 1
        for batch_features, batch_labels in helper.load_preprocess_training_batch(batch_i):
            train_neural_network(sess, optimizer, keep_probability, batch_features, batch_labels)
        print('Epoch {:>2}, CIFAR-10 Batch {}: '.format(epoch + 1, batch_i), end='')
        print_stats(sess, batch_features, batch_labels, cost, accuracy)

```

Checking the Training on a Single Batch...

Epoch 1, CIFAR-10 Batch 1:	Loss: 2.3026	Validation Accuracy: 0.098800
Epoch 2, CIFAR-10 Batch 1:	Loss: 2.3026	Validation Accuracy: 0.099000
Epoch 3, CIFAR-10 Batch 1:	Loss: 2.3026	Validation Accuracy: 0.099000
Epoch 4, CIFAR-10 Batch 1:	Loss: 2.3026	Validation Accuracy: 0.098400
Epoch 5, CIFAR-10 Batch 1:	Loss: 2.3026	Validation Accuracy: 0.098400
Epoch 6, CIFAR-10 Batch 1:	Loss: 2.3026	Validation Accuracy: 0.098400

```
Epoch 7, CIFAR-10 Batch 1: Loss:      2.3026      Validation Accuracy: 0.098400
Epoch 8, CIFAR-10 Batch 1: Loss:      2.3026      Validation Accuracy: 0.098400
Epoch 9, CIFAR-10 Batch 1: Loss:      2.3026      Validation Accuracy: 0.098400
Epoch 10, CIFAR-10 Batch 1: Loss:      2.3026      Validation Accuracy: 0.098200
```

2.2.5 Fully Train the Model

Now that you got a good accuracy with a single CIFAR-10 batch, try it with all five batches.

```
In [192]: """
          DON'T MODIFY ANYTHING IN THIS CELL
          """

          save_model_path = './image_classification'

          print('Training...')
          with tf.Session() as sess:
              # Initializing the variables
              sess.run(tf.global_variables_initializer())

              # Training cycle
              for epoch in range(epochs):
                  # Loop over all batches
                  n_batches = 5
                  for batch_i in range(1, n_batches + 1):
                      for batch_features, batch_labels in helper.load_preprocess_training_batch(
                          train_neural_network(sess, optimizer, keep_probability, batch_features
                      print('Epoch {:>2}, CIFAR-10 Batch {}: '.format(epoch + 1, batch_i), end=
                      print_stats(sess, batch_features, batch_labels, cost, accuracy)

              # Save Model
              saver = tf.train.Saver()
              save_path = saver.save(sess, save_model_path)
```

Training...

```
Epoch 1, CIFAR-10 Batch 1: Loss: 2.2472 Validation Accuracy: 0.1652
Epoch 1, CIFAR-10 Batch 2: Loss: 2.2197 Validation Accuracy: 0.2520
Epoch 1, CIFAR-10 Batch 3: Loss: 1.8955 Validation Accuracy: 0.2932
Epoch 1, CIFAR-10 Batch 4: Loss: 1.8806 Validation Accuracy: 0.3226
Epoch 1, CIFAR-10 Batch 5: Loss: 1.8179 Validation Accuracy: 0.3596
Epoch 2, CIFAR-10 Batch 1: Loss: 1.9865 Validation Accuracy: 0.3766
Epoch 2, CIFAR-10 Batch 2: Loss: 1.7393 Validation Accuracy: 0.4066
Epoch 2, CIFAR-10 Batch 3: Loss: 1.4919 Validation Accuracy: 0.4078
Epoch 2, CIFAR-10 Batch 4: Loss: 1.5969 Validation Accuracy: 0.4258
Epoch 2, CIFAR-10 Batch 5: Loss: 1.6071 Validation Accuracy: 0.4172
Epoch 3, CIFAR-10 Batch 1: Loss: 1.7288 Validation Accuracy: 0.4488
Epoch 3, CIFAR-10 Batch 2: Loss: 1.5289 Validation Accuracy: 0.4710
Epoch 3, CIFAR-10 Batch 3: Loss: 1.2569 Validation Accuracy: 0.4420
```

Epoch 3, CIFAR-10 Batch 4:	Loss: 1.4595	Validation Accuracy: 0.4742
Epoch 3, CIFAR-10 Batch 5:	Loss: 1.4661	Validation Accuracy: 0.4684
Epoch 4, CIFAR-10 Batch 1:	Loss: 1.6106	Validation Accuracy: 0.4840
Epoch 4, CIFAR-10 Batch 2:	Loss: 1.4443	Validation Accuracy: 0.4788
Epoch 4, CIFAR-10 Batch 3:	Loss: 1.1750	Validation Accuracy: 0.4674
Epoch 4, CIFAR-10 Batch 4:	Loss: 1.3708	Validation Accuracy: 0.4972
Epoch 4, CIFAR-10 Batch 5:	Loss: 1.3299	Validation Accuracy: 0.4906
Epoch 5, CIFAR-10 Batch 1:	Loss: 1.5095	Validation Accuracy: 0.5022
Epoch 5, CIFAR-10 Batch 2:	Loss: 1.3289	Validation Accuracy: 0.5068
Epoch 5, CIFAR-10 Batch 3:	Loss: 1.1163	Validation Accuracy: 0.4966
Epoch 5, CIFAR-10 Batch 4:	Loss: 1.3153	Validation Accuracy: 0.5136
Epoch 5, CIFAR-10 Batch 5:	Loss: 1.2742	Validation Accuracy: 0.5222
Epoch 6, CIFAR-10 Batch 1:	Loss: 1.4295	Validation Accuracy: 0.4916
Epoch 6, CIFAR-10 Batch 2:	Loss: 1.2460	Validation Accuracy: 0.5116
Epoch 6, CIFAR-10 Batch 3:	Loss: 1.0543	Validation Accuracy: 0.5184
Epoch 6, CIFAR-10 Batch 4:	Loss: 1.2798	Validation Accuracy: 0.5194
Epoch 6, CIFAR-10 Batch 5:	Loss: 1.2061	Validation Accuracy: 0.5236
Epoch 7, CIFAR-10 Batch 1:	Loss: 1.3290	Validation Accuracy: 0.5198
Epoch 7, CIFAR-10 Batch 2:	Loss: 1.1793	Validation Accuracy: 0.5230
Epoch 7, CIFAR-10 Batch 3:	Loss: 1.0385	Validation Accuracy: 0.5258
Epoch 7, CIFAR-10 Batch 4:	Loss: 1.2688	Validation Accuracy: 0.5262
Epoch 7, CIFAR-10 Batch 5:	Loss: 1.1472	Validation Accuracy: 0.5436
Epoch 8, CIFAR-10 Batch 1:	Loss: 1.2780	Validation Accuracy: 0.5190
Epoch 8, CIFAR-10 Batch 2:	Loss: 1.0998	Validation Accuracy: 0.5406
Epoch 8, CIFAR-10 Batch 3:	Loss: 1.0267	Validation Accuracy: 0.5346
Epoch 8, CIFAR-10 Batch 4:	Loss: 1.2033	Validation Accuracy: 0.5396
Epoch 8, CIFAR-10 Batch 5:	Loss: 1.0755	Validation Accuracy: 0.5492
Epoch 9, CIFAR-10 Batch 1:	Loss: 1.2058	Validation Accuracy: 0.5460
Epoch 9, CIFAR-10 Batch 2:	Loss: 1.0946	Validation Accuracy: 0.5558
Epoch 9, CIFAR-10 Batch 3:	Loss: 0.9586	Validation Accuracy: 0.5426
Epoch 9, CIFAR-10 Batch 4:	Loss: 1.1376	Validation Accuracy: 0.5576
Epoch 9, CIFAR-10 Batch 5:	Loss: 1.0260	Validation Accuracy: 0.5566
Epoch 10, CIFAR-10 Batch 1:	Loss: 1.1731	Validation Accuracy: 0.5382
Epoch 10, CIFAR-10 Batch 2:	Loss: 1.0245	Validation Accuracy: 0.5684
Epoch 10, CIFAR-10 Batch 3:	Loss: 0.9013	Validation Accuracy: 0.5590
Epoch 10, CIFAR-10 Batch 4:	Loss: 1.1070	Validation Accuracy: 0.5620
Epoch 10, CIFAR-10 Batch 5:	Loss: 0.9483	Validation Accuracy: 0.5662
Epoch 11, CIFAR-10 Batch 1:	Loss: 1.0874	Validation Accuracy: 0.5576
Epoch 11, CIFAR-10 Batch 2:	Loss: 0.9886	Validation Accuracy: 0.5732
Epoch 11, CIFAR-10 Batch 3:	Loss: 0.9538	Validation Accuracy: 0.5568
Epoch 11, CIFAR-10 Batch 4:	Loss: 1.0513	Validation Accuracy: 0.5768
Epoch 11, CIFAR-10 Batch 5:	Loss: 0.9476	Validation Accuracy: 0.5678
Epoch 12, CIFAR-10 Batch 1:	Loss: 1.0451	Validation Accuracy: 0.5598
Epoch 12, CIFAR-10 Batch 2:	Loss: 0.9277	Validation Accuracy: 0.5818
Epoch 12, CIFAR-10 Batch 3:	Loss: 0.8347	Validation Accuracy: 0.5772
Epoch 12, CIFAR-10 Batch 4:	Loss: 1.0288	Validation Accuracy: 0.5810
Epoch 12, CIFAR-10 Batch 5:	Loss: 0.8717	Validation Accuracy: 0.5870
Epoch 13, CIFAR-10 Batch 1:	Loss: 1.0185	Validation Accuracy: 0.5616

Epoch 13, CIFAR-10 Batch 2:	Loss: 0.8769	Validation Accuracy: 0.5888
Epoch 13, CIFAR-10 Batch 3:	Loss: 0.8673	Validation Accuracy: 0.5710
Epoch 13, CIFAR-10 Batch 4:	Loss: 0.9661	Validation Accuracy: 0.5870
Epoch 13, CIFAR-10 Batch 5:	Loss: 0.8071	Validation Accuracy: 0.5940
Epoch 14, CIFAR-10 Batch 1:	Loss: 0.9513	Validation Accuracy: 0.5686
Epoch 14, CIFAR-10 Batch 2:	Loss: 0.8549	Validation Accuracy: 0.5882
Epoch 14, CIFAR-10 Batch 3:	Loss: 0.7269	Validation Accuracy: 0.5890
Epoch 14, CIFAR-10 Batch 4:	Loss: 0.9111	Validation Accuracy: 0.5958
Epoch 14, CIFAR-10 Batch 5:	Loss: 0.7728	Validation Accuracy: 0.6014
Epoch 15, CIFAR-10 Batch 1:	Loss: 0.9081	Validation Accuracy: 0.5788
Epoch 15, CIFAR-10 Batch 2:	Loss: 0.7984	Validation Accuracy: 0.5886
Epoch 15, CIFAR-10 Batch 3:	Loss: 0.7086	Validation Accuracy: 0.5924
Epoch 15, CIFAR-10 Batch 4:	Loss: 0.9115	Validation Accuracy: 0.5930
Epoch 15, CIFAR-10 Batch 5:	Loss: 0.7204	Validation Accuracy: 0.6058
Epoch 16, CIFAR-10 Batch 1:	Loss: 0.9075	Validation Accuracy: 0.5750
Epoch 16, CIFAR-10 Batch 2:	Loss: 0.7984	Validation Accuracy: 0.5966
Epoch 16, CIFAR-10 Batch 3:	Loss: 0.6322	Validation Accuracy: 0.6044
Epoch 16, CIFAR-10 Batch 4:	Loss: 0.8999	Validation Accuracy: 0.5962
Epoch 16, CIFAR-10 Batch 5:	Loss: 0.6865	Validation Accuracy: 0.6122
Epoch 17, CIFAR-10 Batch 1:	Loss: 0.8466	Validation Accuracy: 0.5810
Epoch 17, CIFAR-10 Batch 2:	Loss: 0.7507	Validation Accuracy: 0.5976
Epoch 17, CIFAR-10 Batch 3:	Loss: 0.6670	Validation Accuracy: 0.5914
Epoch 17, CIFAR-10 Batch 4:	Loss: 0.7901	Validation Accuracy: 0.6078
Epoch 17, CIFAR-10 Batch 5:	Loss: 0.6369	Validation Accuracy: 0.6048
Epoch 18, CIFAR-10 Batch 1:	Loss: 0.8253	Validation Accuracy: 0.5966
Epoch 18, CIFAR-10 Batch 2:	Loss: 0.7154	Validation Accuracy: 0.6024
Epoch 18, CIFAR-10 Batch 3:	Loss: 0.6291	Validation Accuracy: 0.6072
Epoch 18, CIFAR-10 Batch 4:	Loss: 0.8043	Validation Accuracy: 0.6096
Epoch 18, CIFAR-10 Batch 5:	Loss: 0.6249	Validation Accuracy: 0.6148
Epoch 19, CIFAR-10 Batch 1:	Loss: 0.7903	Validation Accuracy: 0.6060
Epoch 19, CIFAR-10 Batch 2:	Loss: 0.7113	Validation Accuracy: 0.6032
Epoch 19, CIFAR-10 Batch 3:	Loss: 0.5891	Validation Accuracy: 0.6060
Epoch 19, CIFAR-10 Batch 4:	Loss: 0.7701	Validation Accuracy: 0.6116
Epoch 19, CIFAR-10 Batch 5:	Loss: 0.5846	Validation Accuracy: 0.6148
Epoch 20, CIFAR-10 Batch 1:	Loss: 0.7700	Validation Accuracy: 0.6118
Epoch 20, CIFAR-10 Batch 2:	Loss: 0.6895	Validation Accuracy: 0.6082
Epoch 20, CIFAR-10 Batch 3:	Loss: 0.5449	Validation Accuracy: 0.6106
Epoch 20, CIFAR-10 Batch 4:	Loss: 0.7440	Validation Accuracy: 0.6232
Epoch 20, CIFAR-10 Batch 5:	Loss: 0.5239	Validation Accuracy: 0.6306
Epoch 21, CIFAR-10 Batch 1:	Loss: 0.7096	Validation Accuracy: 0.6192
Epoch 21, CIFAR-10 Batch 2:	Loss: 0.6403	Validation Accuracy: 0.6226
Epoch 21, CIFAR-10 Batch 3:	Loss: 0.5562	Validation Accuracy: 0.6124
Epoch 21, CIFAR-10 Batch 4:	Loss: 0.6908	Validation Accuracy: 0.6234
Epoch 21, CIFAR-10 Batch 5:	Loss: 0.5232	Validation Accuracy: 0.6326
Epoch 22, CIFAR-10 Batch 1:	Loss: 0.6805	Validation Accuracy: 0.6148
Epoch 22, CIFAR-10 Batch 2:	Loss: 0.6282	Validation Accuracy: 0.6044
Epoch 22, CIFAR-10 Batch 3:	Loss: 0.5308	Validation Accuracy: 0.6270
Epoch 22, CIFAR-10 Batch 4:	Loss: 0.6384	Validation Accuracy: 0.6316

Epoch 22, CIFAR-10 Batch 5:	Loss: 0.4973	Validation Accuracy: 0.6316
Epoch 23, CIFAR-10 Batch 1:	Loss: 0.6599	Validation Accuracy: 0.6086
Epoch 23, CIFAR-10 Batch 2:	Loss: 0.5832	Validation Accuracy: 0.6258
Epoch 23, CIFAR-10 Batch 3:	Loss: 0.5687	Validation Accuracy: 0.6170
Epoch 23, CIFAR-10 Batch 4:	Loss: 0.6817	Validation Accuracy: 0.6182
Epoch 23, CIFAR-10 Batch 5:	Loss: 0.4585	Validation Accuracy: 0.6320
Epoch 24, CIFAR-10 Batch 1:	Loss: 0.6522	Validation Accuracy: 0.6106
Epoch 24, CIFAR-10 Batch 2:	Loss: 0.6002	Validation Accuracy: 0.6052
Epoch 24, CIFAR-10 Batch 3:	Loss: 0.5057	Validation Accuracy: 0.6228
Epoch 24, CIFAR-10 Batch 4:	Loss: 0.6265	Validation Accuracy: 0.6254
Epoch 24, CIFAR-10 Batch 5:	Loss: 0.4791	Validation Accuracy: 0.6298
Epoch 25, CIFAR-10 Batch 1:	Loss: 0.6356	Validation Accuracy: 0.6148
Epoch 25, CIFAR-10 Batch 2:	Loss: 0.5806	Validation Accuracy: 0.6328
Epoch 25, CIFAR-10 Batch 3:	Loss: 0.5068	Validation Accuracy: 0.6184
Epoch 25, CIFAR-10 Batch 4:	Loss: 0.5947	Validation Accuracy: 0.6224
Epoch 25, CIFAR-10 Batch 5:	Loss: 0.4412	Validation Accuracy: 0.6322
Epoch 26, CIFAR-10 Batch 1:	Loss: 0.6127	Validation Accuracy: 0.6096
Epoch 26, CIFAR-10 Batch 2:	Loss: 0.5621	Validation Accuracy: 0.6232
Epoch 26, CIFAR-10 Batch 3:	Loss: 0.4865	Validation Accuracy: 0.6114
Epoch 26, CIFAR-10 Batch 4:	Loss: 0.6060	Validation Accuracy: 0.6350
Epoch 26, CIFAR-10 Batch 5:	Loss: 0.4132	Validation Accuracy: 0.6396
Epoch 27, CIFAR-10 Batch 1:	Loss: 0.6192	Validation Accuracy: 0.6164
Epoch 27, CIFAR-10 Batch 2:	Loss: 0.5595	Validation Accuracy: 0.6350
Epoch 27, CIFAR-10 Batch 3:	Loss: 0.4382	Validation Accuracy: 0.6314
Epoch 27, CIFAR-10 Batch 4:	Loss: 0.6016	Validation Accuracy: 0.6298
Epoch 27, CIFAR-10 Batch 5:	Loss: 0.4045	Validation Accuracy: 0.6426
Epoch 28, CIFAR-10 Batch 1:	Loss: 0.6030	Validation Accuracy: 0.6110
Epoch 28, CIFAR-10 Batch 2:	Loss: 0.4863	Validation Accuracy: 0.6368
Epoch 28, CIFAR-10 Batch 3:	Loss: 0.4247	Validation Accuracy: 0.6394
Epoch 28, CIFAR-10 Batch 4:	Loss: 0.5546	Validation Accuracy: 0.6358
Epoch 28, CIFAR-10 Batch 5:	Loss: 0.3845	Validation Accuracy: 0.6372
Epoch 29, CIFAR-10 Batch 1:	Loss: 0.5878	Validation Accuracy: 0.6282
Epoch 29, CIFAR-10 Batch 2:	Loss: 0.4981	Validation Accuracy: 0.6434
Epoch 29, CIFAR-10 Batch 3:	Loss: 0.4222	Validation Accuracy: 0.6200
Epoch 29, CIFAR-10 Batch 4:	Loss: 0.5678	Validation Accuracy: 0.6290
Epoch 29, CIFAR-10 Batch 5:	Loss: 0.3806	Validation Accuracy: 0.6408
Epoch 30, CIFAR-10 Batch 1:	Loss: 0.5599	Validation Accuracy: 0.6292
Epoch 30, CIFAR-10 Batch 2:	Loss: 0.4545	Validation Accuracy: 0.6416
Epoch 30, CIFAR-10 Batch 3:	Loss: 0.3705	Validation Accuracy: 0.6382
Epoch 30, CIFAR-10 Batch 4:	Loss: 0.5332	Validation Accuracy: 0.6250
Epoch 30, CIFAR-10 Batch 5:	Loss: 0.3756	Validation Accuracy: 0.6470
Epoch 31, CIFAR-10 Batch 1:	Loss: 0.5463	Validation Accuracy: 0.6372
Epoch 31, CIFAR-10 Batch 2:	Loss: 0.4421	Validation Accuracy: 0.6290
Epoch 31, CIFAR-10 Batch 3:	Loss: 0.3968	Validation Accuracy: 0.6372
Epoch 31, CIFAR-10 Batch 4:	Loss: 0.5056	Validation Accuracy: 0.6358
Epoch 31, CIFAR-10 Batch 5:	Loss: 0.3685	Validation Accuracy: 0.6388
Epoch 32, CIFAR-10 Batch 1:	Loss: 0.5519	Validation Accuracy: 0.6378
Epoch 32, CIFAR-10 Batch 2:	Loss: 0.4468	Validation Accuracy: 0.6344

Epoch 32, CIFAR-10 Batch 3:	Loss: 0.3848	Validation Accuracy: 0.6244
Epoch 32, CIFAR-10 Batch 4:	Loss: 0.4847	Validation Accuracy: 0.6336
Epoch 32, CIFAR-10 Batch 5:	Loss: 0.3616	Validation Accuracy: 0.6462
Epoch 33, CIFAR-10 Batch 1:	Loss: 0.5571	Validation Accuracy: 0.6188
Epoch 33, CIFAR-10 Batch 2:	Loss: 0.4692	Validation Accuracy: 0.6280
Epoch 33, CIFAR-10 Batch 3:	Loss: 0.3964	Validation Accuracy: 0.6068
Epoch 33, CIFAR-10 Batch 4:	Loss: 0.4881	Validation Accuracy: 0.6304
Epoch 33, CIFAR-10 Batch 5:	Loss: 0.3554	Validation Accuracy: 0.6466
Epoch 34, CIFAR-10 Batch 1:	Loss: 0.4817	Validation Accuracy: 0.6322
Epoch 34, CIFAR-10 Batch 2:	Loss: 0.4474	Validation Accuracy: 0.6322
Epoch 34, CIFAR-10 Batch 3:	Loss: 0.3523	Validation Accuracy: 0.6366
Epoch 34, CIFAR-10 Batch 4:	Loss: 0.4233	Validation Accuracy: 0.6356
Epoch 34, CIFAR-10 Batch 5:	Loss: 0.3350	Validation Accuracy: 0.6522
Epoch 35, CIFAR-10 Batch 1:	Loss: 0.4697	Validation Accuracy: 0.6370
Epoch 35, CIFAR-10 Batch 2:	Loss: 0.4731	Validation Accuracy: 0.6316
Epoch 35, CIFAR-10 Batch 3:	Loss: 0.3371	Validation Accuracy: 0.6344
Epoch 35, CIFAR-10 Batch 4:	Loss: 0.5013	Validation Accuracy: 0.6062
Epoch 35, CIFAR-10 Batch 5:	Loss: 0.3222	Validation Accuracy: 0.6386
Epoch 36, CIFAR-10 Batch 1:	Loss: 0.4564	Validation Accuracy: 0.6292
Epoch 36, CIFAR-10 Batch 2:	Loss: 0.4307	Validation Accuracy: 0.6324
Epoch 36, CIFAR-10 Batch 3:	Loss: 0.3396	Validation Accuracy: 0.6132
Epoch 36, CIFAR-10 Batch 4:	Loss: 0.4674	Validation Accuracy: 0.6226
Epoch 36, CIFAR-10 Batch 5:	Loss: 0.3108	Validation Accuracy: 0.6438
Epoch 37, CIFAR-10 Batch 1:	Loss: 0.4615	Validation Accuracy: 0.6274
Epoch 37, CIFAR-10 Batch 2:	Loss: 0.4311	Validation Accuracy: 0.6306
Epoch 37, CIFAR-10 Batch 3:	Loss: 0.3537	Validation Accuracy: 0.6186
Epoch 37, CIFAR-10 Batch 4:	Loss: 0.4242	Validation Accuracy: 0.6342
Epoch 37, CIFAR-10 Batch 5:	Loss: 0.3075	Validation Accuracy: 0.6486
Epoch 38, CIFAR-10 Batch 1:	Loss: 0.4387	Validation Accuracy: 0.6364
Epoch 38, CIFAR-10 Batch 2:	Loss: 0.4072	Validation Accuracy: 0.6364
Epoch 38, CIFAR-10 Batch 3:	Loss: 0.3738	Validation Accuracy: 0.5998
Epoch 38, CIFAR-10 Batch 4:	Loss: 0.4240	Validation Accuracy: 0.6334
Epoch 38, CIFAR-10 Batch 5:	Loss: 0.3304	Validation Accuracy: 0.6458
Epoch 39, CIFAR-10 Batch 1:	Loss: 0.4379	Validation Accuracy: 0.6418
Epoch 39, CIFAR-10 Batch 2:	Loss: 0.3932	Validation Accuracy: 0.6472
Epoch 39, CIFAR-10 Batch 3:	Loss: 0.3375	Validation Accuracy: 0.6394
Epoch 39, CIFAR-10 Batch 4:	Loss: 0.4067	Validation Accuracy: 0.6474
Epoch 39, CIFAR-10 Batch 5:	Loss: 0.2993	Validation Accuracy: 0.6454
Epoch 40, CIFAR-10 Batch 1:	Loss: 0.4051	Validation Accuracy: 0.6432
Epoch 40, CIFAR-10 Batch 2:	Loss: 0.3624	Validation Accuracy: 0.6428
Epoch 40, CIFAR-10 Batch 3:	Loss: 0.3075	Validation Accuracy: 0.6276
Epoch 40, CIFAR-10 Batch 4:	Loss: 0.3944	Validation Accuracy: 0.6370
Epoch 40, CIFAR-10 Batch 5:	Loss: 0.2922	Validation Accuracy: 0.6486
Epoch 41, CIFAR-10 Batch 1:	Loss: 0.4176	Validation Accuracy: 0.6352
Epoch 41, CIFAR-10 Batch 2:	Loss: 0.3518	Validation Accuracy: 0.6526
Epoch 41, CIFAR-10 Batch 3:	Loss: 0.3209	Validation Accuracy: 0.6462
Epoch 41, CIFAR-10 Batch 4:	Loss: 0.3911	Validation Accuracy: 0.6422
Epoch 41, CIFAR-10 Batch 5:	Loss: 0.3045	Validation Accuracy: 0.6450

Epoch 42, CIFAR-10 Batch 1:	Loss: 0.3912	Validation Accuracy: 0.6416
Epoch 42, CIFAR-10 Batch 2:	Loss: 0.3448	Validation Accuracy: 0.6500
Epoch 42, CIFAR-10 Batch 3:	Loss: 0.2905	Validation Accuracy: 0.6374
Epoch 42, CIFAR-10 Batch 4:	Loss: 0.3845	Validation Accuracy: 0.6428
Epoch 42, CIFAR-10 Batch 5:	Loss: 0.2831	Validation Accuracy: 0.6468
Epoch 43, CIFAR-10 Batch 1:	Loss: 0.3866	Validation Accuracy: 0.6466
Epoch 43, CIFAR-10 Batch 2:	Loss: 0.3314	Validation Accuracy: 0.6296
Epoch 43, CIFAR-10 Batch 3:	Loss: 0.3027	Validation Accuracy: 0.6252
Epoch 43, CIFAR-10 Batch 4:	Loss: 0.3716	Validation Accuracy: 0.6330
Epoch 43, CIFAR-10 Batch 5:	Loss: 0.2905	Validation Accuracy: 0.6324
Epoch 44, CIFAR-10 Batch 1:	Loss: 0.3709	Validation Accuracy: 0.6348
Epoch 44, CIFAR-10 Batch 2:	Loss: 0.3476	Validation Accuracy: 0.6312
Epoch 44, CIFAR-10 Batch 3:	Loss: 0.2899	Validation Accuracy: 0.6252
Epoch 44, CIFAR-10 Batch 4:	Loss: 0.3553	Validation Accuracy: 0.6296
Epoch 44, CIFAR-10 Batch 5:	Loss: 0.2717	Validation Accuracy: 0.6392
Epoch 45, CIFAR-10 Batch 1:	Loss: 0.3694	Validation Accuracy: 0.6360
Epoch 45, CIFAR-10 Batch 2:	Loss: 0.3413	Validation Accuracy: 0.6304
Epoch 45, CIFAR-10 Batch 3:	Loss: 0.2766	Validation Accuracy: 0.6414
Epoch 45, CIFAR-10 Batch 4:	Loss: 0.3524	Validation Accuracy: 0.6304
Epoch 45, CIFAR-10 Batch 5:	Loss: 0.2908	Validation Accuracy: 0.6438
Epoch 46, CIFAR-10 Batch 1:	Loss: 0.3792	Validation Accuracy: 0.6404
Epoch 46, CIFAR-10 Batch 2:	Loss: 0.3384	Validation Accuracy: 0.6344
Epoch 46, CIFAR-10 Batch 3:	Loss: 0.2679	Validation Accuracy: 0.6324
Epoch 46, CIFAR-10 Batch 4:	Loss: 0.3421	Validation Accuracy: 0.6396
Epoch 46, CIFAR-10 Batch 5:	Loss: 0.2892	Validation Accuracy: 0.6428
Epoch 47, CIFAR-10 Batch 1:	Loss: 0.3916	Validation Accuracy: 0.6462
Epoch 47, CIFAR-10 Batch 2:	Loss: 0.3372	Validation Accuracy: 0.6394
Epoch 47, CIFAR-10 Batch 3:	Loss: 0.2645	Validation Accuracy: 0.6396
Epoch 47, CIFAR-10 Batch 4:	Loss: 0.3455	Validation Accuracy: 0.6540
Epoch 47, CIFAR-10 Batch 5:	Loss: 0.2777	Validation Accuracy: 0.6320
Epoch 48, CIFAR-10 Batch 1:	Loss: 0.3612	Validation Accuracy: 0.6544
Epoch 48, CIFAR-10 Batch 2:	Loss: 0.3094	Validation Accuracy: 0.6454
Epoch 48, CIFAR-10 Batch 3:	Loss: 0.2499	Validation Accuracy: 0.6562
Epoch 48, CIFAR-10 Batch 4:	Loss: 0.3579	Validation Accuracy: 0.6436
Epoch 48, CIFAR-10 Batch 5:	Loss: 0.2731	Validation Accuracy: 0.6492
Epoch 49, CIFAR-10 Batch 1:	Loss: 0.3780	Validation Accuracy: 0.6418
Epoch 49, CIFAR-10 Batch 2:	Loss: 0.3005	Validation Accuracy: 0.6324
Epoch 49, CIFAR-10 Batch 3:	Loss: 0.2516	Validation Accuracy: 0.6512
Epoch 49, CIFAR-10 Batch 4:	Loss: 0.3675	Validation Accuracy: 0.6340
Epoch 49, CIFAR-10 Batch 5:	Loss: 0.2883	Validation Accuracy: 0.6420
Epoch 50, CIFAR-10 Batch 1:	Loss: 0.4280	Validation Accuracy: 0.6308
Epoch 50, CIFAR-10 Batch 2:	Loss: 0.3101	Validation Accuracy: 0.6422
Epoch 50, CIFAR-10 Batch 3:	Loss: 0.2526	Validation Accuracy: 0.6554
Epoch 50, CIFAR-10 Batch 4:	Loss: 0.3487	Validation Accuracy: 0.6348
Epoch 50, CIFAR-10 Batch 5:	Loss: 0.2617	Validation Accuracy: 0.6418
Epoch 51, CIFAR-10 Batch 1:	Loss: 0.3641	Validation Accuracy: 0.6394
Epoch 51, CIFAR-10 Batch 2:	Loss: 0.2870	Validation Accuracy: 0.6398
Epoch 51, CIFAR-10 Batch 3:	Loss: 0.2411	Validation Accuracy: 0.6546

Epoch 51, CIFAR-10 Batch 4:	Loss: 0.3327	Validation Accuracy: 0.6348
Epoch 51, CIFAR-10 Batch 5:	Loss: 0.2501	Validation Accuracy: 0.6488
Epoch 52, CIFAR-10 Batch 1:	Loss: 0.3731	Validation Accuracy: 0.6390
Epoch 52, CIFAR-10 Batch 2:	Loss: 0.2928	Validation Accuracy: 0.6406
Epoch 52, CIFAR-10 Batch 3:	Loss: 0.2405	Validation Accuracy: 0.6568
Epoch 52, CIFAR-10 Batch 4:	Loss: 0.3327	Validation Accuracy: 0.6334
Epoch 52, CIFAR-10 Batch 5:	Loss: 0.2160	Validation Accuracy: 0.6474
Epoch 53, CIFAR-10 Batch 1:	Loss: 0.3493	Validation Accuracy: 0.6344
Epoch 53, CIFAR-10 Batch 2:	Loss: 0.2666	Validation Accuracy: 0.6428
Epoch 53, CIFAR-10 Batch 3:	Loss: 0.2388	Validation Accuracy: 0.6454
Epoch 53, CIFAR-10 Batch 4:	Loss: 0.3200	Validation Accuracy: 0.6378
Epoch 53, CIFAR-10 Batch 5:	Loss: 0.2300	Validation Accuracy: 0.6528
Epoch 54, CIFAR-10 Batch 1:	Loss: 0.3733	Validation Accuracy: 0.6178
Epoch 54, CIFAR-10 Batch 2:	Loss: 0.3019	Validation Accuracy: 0.6290
Epoch 54, CIFAR-10 Batch 3:	Loss: 0.2400	Validation Accuracy: 0.6376
Epoch 54, CIFAR-10 Batch 4:	Loss: 0.3577	Validation Accuracy: 0.6276
Epoch 54, CIFAR-10 Batch 5:	Loss: 0.2236	Validation Accuracy: 0.6530
Epoch 55, CIFAR-10 Batch 1:	Loss: 0.3323	Validation Accuracy: 0.6360
Epoch 55, CIFAR-10 Batch 2:	Loss: 0.2600	Validation Accuracy: 0.6308
Epoch 55, CIFAR-10 Batch 3:	Loss: 0.2206	Validation Accuracy: 0.6520
Epoch 55, CIFAR-10 Batch 4:	Loss: 0.3429	Validation Accuracy: 0.6312
Epoch 55, CIFAR-10 Batch 5:	Loss: 0.2190	Validation Accuracy: 0.6502
Epoch 56, CIFAR-10 Batch 1:	Loss: 0.3055	Validation Accuracy: 0.6400
Epoch 56, CIFAR-10 Batch 2:	Loss: 0.2594	Validation Accuracy: 0.6422
Epoch 56, CIFAR-10 Batch 3:	Loss: 0.2274	Validation Accuracy: 0.6526
Epoch 56, CIFAR-10 Batch 4:	Loss: 0.3353	Validation Accuracy: 0.6284
Epoch 56, CIFAR-10 Batch 5:	Loss: 0.2343	Validation Accuracy: 0.6504
Epoch 57, CIFAR-10 Batch 1:	Loss: 0.3137	Validation Accuracy: 0.6350
Epoch 57, CIFAR-10 Batch 2:	Loss: 0.2665	Validation Accuracy: 0.6426
Epoch 57, CIFAR-10 Batch 3:	Loss: 0.2312	Validation Accuracy: 0.6540
Epoch 57, CIFAR-10 Batch 4:	Loss: 0.2990	Validation Accuracy: 0.6476
Epoch 57, CIFAR-10 Batch 5:	Loss: 0.2097	Validation Accuracy: 0.6534
Epoch 58, CIFAR-10 Batch 1:	Loss: 0.3075	Validation Accuracy: 0.6314
Epoch 58, CIFAR-10 Batch 2:	Loss: 0.2558	Validation Accuracy: 0.6470
Epoch 58, CIFAR-10 Batch 3:	Loss: 0.2645	Validation Accuracy: 0.6482
Epoch 58, CIFAR-10 Batch 4:	Loss: 0.3015	Validation Accuracy: 0.6478
Epoch 58, CIFAR-10 Batch 5:	Loss: 0.2121	Validation Accuracy: 0.6438
Epoch 59, CIFAR-10 Batch 1:	Loss: 0.3046	Validation Accuracy: 0.6440
Epoch 59, CIFAR-10 Batch 2:	Loss: 0.2689	Validation Accuracy: 0.6296
Epoch 59, CIFAR-10 Batch 3:	Loss: 0.2326	Validation Accuracy: 0.6476
Epoch 59, CIFAR-10 Batch 4:	Loss: 0.3182	Validation Accuracy: 0.6344
Epoch 59, CIFAR-10 Batch 5:	Loss: 0.2111	Validation Accuracy: 0.6408
Epoch 60, CIFAR-10 Batch 1:	Loss: 0.3246	Validation Accuracy: 0.6378
Epoch 60, CIFAR-10 Batch 2:	Loss: 0.2648	Validation Accuracy: 0.6424
Epoch 60, CIFAR-10 Batch 3:	Loss: 0.2256	Validation Accuracy: 0.6538
Epoch 60, CIFAR-10 Batch 4:	Loss: 0.3239	Validation Accuracy: 0.6210
Epoch 60, CIFAR-10 Batch 5:	Loss: 0.2322	Validation Accuracy: 0.6274
Epoch 61, CIFAR-10 Batch 1:	Loss: 0.2865	Validation Accuracy: 0.6496

Epoch 61, CIFAR-10 Batch 2:	Loss: 0.2661	Validation Accuracy: 0.6438
Epoch 61, CIFAR-10 Batch 3:	Loss: 0.2210	Validation Accuracy: 0.6512
Epoch 61, CIFAR-10 Batch 4:	Loss: 0.3262	Validation Accuracy: 0.6090
Epoch 61, CIFAR-10 Batch 5:	Loss: 0.2232	Validation Accuracy: 0.6228
Epoch 62, CIFAR-10 Batch 1:	Loss: 0.2834	Validation Accuracy: 0.6456
Epoch 62, CIFAR-10 Batch 2:	Loss: 0.2854	Validation Accuracy: 0.6276
Epoch 62, CIFAR-10 Batch 3:	Loss: 0.2391	Validation Accuracy: 0.6516
Epoch 62, CIFAR-10 Batch 4:	Loss: 0.3325	Validation Accuracy: 0.6140
Epoch 62, CIFAR-10 Batch 5:	Loss: 0.2150	Validation Accuracy: 0.6402
Epoch 63, CIFAR-10 Batch 1:	Loss: 0.2650	Validation Accuracy: 0.6428
Epoch 63, CIFAR-10 Batch 2:	Loss: 0.2696	Validation Accuracy: 0.6236
Epoch 63, CIFAR-10 Batch 3:	Loss: 0.2413	Validation Accuracy: 0.6490
Epoch 63, CIFAR-10 Batch 4:	Loss: 0.2805	Validation Accuracy: 0.6308
Epoch 63, CIFAR-10 Batch 5:	Loss: 0.2113	Validation Accuracy: 0.6400
Epoch 64, CIFAR-10 Batch 1:	Loss: 0.2712	Validation Accuracy: 0.6426
Epoch 64, CIFAR-10 Batch 2:	Loss: 0.2723	Validation Accuracy: 0.6206
Epoch 64, CIFAR-10 Batch 3:	Loss: 0.2622	Validation Accuracy: 0.6412
Epoch 64, CIFAR-10 Batch 4:	Loss: 0.3032	Validation Accuracy: 0.6282
Epoch 64, CIFAR-10 Batch 5:	Loss: 0.2025	Validation Accuracy: 0.6464
Epoch 65, CIFAR-10 Batch 1:	Loss: 0.2795	Validation Accuracy: 0.6376
Epoch 65, CIFAR-10 Batch 2:	Loss: 0.2489	Validation Accuracy: 0.6264
Epoch 65, CIFAR-10 Batch 3:	Loss: 0.2479	Validation Accuracy: 0.6322
Epoch 65, CIFAR-10 Batch 4:	Loss: 0.2795	Validation Accuracy: 0.6378
Epoch 65, CIFAR-10 Batch 5:	Loss: 0.2124	Validation Accuracy: 0.6358
Epoch 66, CIFAR-10 Batch 1:	Loss: 0.2587	Validation Accuracy: 0.6452
Epoch 66, CIFAR-10 Batch 2:	Loss: 0.2988	Validation Accuracy: 0.6108
Epoch 66, CIFAR-10 Batch 3:	Loss: 0.2073	Validation Accuracy: 0.6386
Epoch 66, CIFAR-10 Batch 4:	Loss: 0.2663	Validation Accuracy: 0.6330
Epoch 66, CIFAR-10 Batch 5:	Loss: 0.2007	Validation Accuracy: 0.6486
Epoch 67, CIFAR-10 Batch 1:	Loss: 0.2546	Validation Accuracy: 0.6434
Epoch 67, CIFAR-10 Batch 2:	Loss: 0.2880	Validation Accuracy: 0.6150
Epoch 67, CIFAR-10 Batch 3:	Loss: 0.2007	Validation Accuracy: 0.6262
Epoch 67, CIFAR-10 Batch 4:	Loss: 0.2541	Validation Accuracy: 0.6290
Epoch 67, CIFAR-10 Batch 5:	Loss: 0.2021	Validation Accuracy: 0.6390
Epoch 68, CIFAR-10 Batch 1:	Loss: 0.2647	Validation Accuracy: 0.6356
Epoch 68, CIFAR-10 Batch 2:	Loss: 0.2966	Validation Accuracy: 0.6036
Epoch 68, CIFAR-10 Batch 3:	Loss: 0.2285	Validation Accuracy: 0.6198
Epoch 68, CIFAR-10 Batch 4:	Loss: 0.2439	Validation Accuracy: 0.6390
Epoch 68, CIFAR-10 Batch 5:	Loss: 0.2083	Validation Accuracy: 0.6400
Epoch 69, CIFAR-10 Batch 1:	Loss: 0.2751	Validation Accuracy: 0.6436
Epoch 69, CIFAR-10 Batch 2:	Loss: 0.2694	Validation Accuracy: 0.6266
Epoch 69, CIFAR-10 Batch 3:	Loss: 0.2105	Validation Accuracy: 0.6280
Epoch 69, CIFAR-10 Batch 4:	Loss: 0.2398	Validation Accuracy: 0.6422
Epoch 69, CIFAR-10 Batch 5:	Loss: 0.1950	Validation Accuracy: 0.6372
Epoch 70, CIFAR-10 Batch 1:	Loss: 0.2716	Validation Accuracy: 0.6420
Epoch 70, CIFAR-10 Batch 2:	Loss: 0.2436	Validation Accuracy: 0.6402
Epoch 70, CIFAR-10 Batch 3:	Loss: 0.2004	Validation Accuracy: 0.6508
Epoch 70, CIFAR-10 Batch 4:	Loss: 0.2300	Validation Accuracy: 0.6410

Epoch 70, CIFAR-10 Batch 5:	Loss: 0.1930	Validation Accuracy: 0.6328
Epoch 71, CIFAR-10 Batch 1:	Loss: 0.2587	Validation Accuracy: 0.6426
Epoch 71, CIFAR-10 Batch 2:	Loss: 0.2366	Validation Accuracy: 0.6370
Epoch 71, CIFAR-10 Batch 3:	Loss: 0.1911	Validation Accuracy: 0.6486
Epoch 71, CIFAR-10 Batch 4:	Loss: 0.2198	Validation Accuracy: 0.6358
Epoch 71, CIFAR-10 Batch 5:	Loss: 0.1809	Validation Accuracy: 0.6346
Epoch 72, CIFAR-10 Batch 1:	Loss: 0.2598	Validation Accuracy: 0.6372
Epoch 72, CIFAR-10 Batch 2:	Loss: 0.2436	Validation Accuracy: 0.6452
Epoch 72, CIFAR-10 Batch 3:	Loss: 0.2185	Validation Accuracy: 0.6304
Epoch 72, CIFAR-10 Batch 4:	Loss: 0.2427	Validation Accuracy: 0.6314
Epoch 72, CIFAR-10 Batch 5:	Loss: 0.1732	Validation Accuracy: 0.6368
Epoch 73, CIFAR-10 Batch 1:	Loss: 0.2552	Validation Accuracy: 0.6416
Epoch 73, CIFAR-10 Batch 2:	Loss: 0.2245	Validation Accuracy: 0.6466
Epoch 73, CIFAR-10 Batch 3:	Loss: 0.1961	Validation Accuracy: 0.6438
Epoch 73, CIFAR-10 Batch 4:	Loss: 0.2209	Validation Accuracy: 0.6416
Epoch 73, CIFAR-10 Batch 5:	Loss: 0.1814	Validation Accuracy: 0.6326
Epoch 74, CIFAR-10 Batch 1:	Loss: 0.2662	Validation Accuracy: 0.6360
Epoch 74, CIFAR-10 Batch 2:	Loss: 0.2216	Validation Accuracy: 0.6350
Epoch 74, CIFAR-10 Batch 3:	Loss: 0.1966	Validation Accuracy: 0.6484
Epoch 74, CIFAR-10 Batch 4:	Loss: 0.2557	Validation Accuracy: 0.6176
Epoch 74, CIFAR-10 Batch 5:	Loss: 0.1733	Validation Accuracy: 0.6292
Epoch 75, CIFAR-10 Batch 1:	Loss: 0.2501	Validation Accuracy: 0.6418
Epoch 75, CIFAR-10 Batch 2:	Loss: 0.2571	Validation Accuracy: 0.6328
Epoch 75, CIFAR-10 Batch 3:	Loss: 0.2178	Validation Accuracy: 0.6292
Epoch 75, CIFAR-10 Batch 4:	Loss: 0.2362	Validation Accuracy: 0.6288
Epoch 75, CIFAR-10 Batch 5:	Loss: 0.1634	Validation Accuracy: 0.6386
Epoch 76, CIFAR-10 Batch 1:	Loss: 0.2718	Validation Accuracy: 0.6322
Epoch 76, CIFAR-10 Batch 2:	Loss: 0.2238	Validation Accuracy: 0.6388
Epoch 76, CIFAR-10 Batch 3:	Loss: 0.1720	Validation Accuracy: 0.6512
Epoch 76, CIFAR-10 Batch 4:	Loss: 0.2498	Validation Accuracy: 0.6192
Epoch 76, CIFAR-10 Batch 5:	Loss: 0.1589	Validation Accuracy: 0.6454
Epoch 77, CIFAR-10 Batch 1:	Loss: 0.2742	Validation Accuracy: 0.6272
Epoch 77, CIFAR-10 Batch 2:	Loss: 0.2427	Validation Accuracy: 0.6346
Epoch 77, CIFAR-10 Batch 3:	Loss: 0.1961	Validation Accuracy: 0.6418
Epoch 77, CIFAR-10 Batch 4:	Loss: 0.2515	Validation Accuracy: 0.6248
Epoch 77, CIFAR-10 Batch 5:	Loss: 0.1556	Validation Accuracy: 0.6414
Epoch 78, CIFAR-10 Batch 1:	Loss: 0.2605	Validation Accuracy: 0.6272
Epoch 78, CIFAR-10 Batch 2:	Loss: 0.2515	Validation Accuracy: 0.6420
Epoch 78, CIFAR-10 Batch 3:	Loss: 0.1898	Validation Accuracy: 0.6500
Epoch 78, CIFAR-10 Batch 4:	Loss: 0.2189	Validation Accuracy: 0.6356
Epoch 78, CIFAR-10 Batch 5:	Loss: 0.1661	Validation Accuracy: 0.6428
Epoch 79, CIFAR-10 Batch 1:	Loss: 0.2716	Validation Accuracy: 0.6246
Epoch 79, CIFAR-10 Batch 2:	Loss: 0.2443	Validation Accuracy: 0.6450
Epoch 79, CIFAR-10 Batch 3:	Loss: 0.1784	Validation Accuracy: 0.6496
Epoch 79, CIFAR-10 Batch 4:	Loss: 0.2389	Validation Accuracy: 0.6320
Epoch 79, CIFAR-10 Batch 5:	Loss: 0.1790	Validation Accuracy: 0.6336
Epoch 80, CIFAR-10 Batch 1:	Loss: 0.3127	Validation Accuracy: 0.5988
Epoch 80, CIFAR-10 Batch 2:	Loss: 0.2769	Validation Accuracy: 0.6244

Epoch 80, CIFAR-10 Batch 3:	Loss: 0.1831	Validation Accuracy: 0.6418
Epoch 80, CIFAR-10 Batch 4:	Loss: 0.2196	Validation Accuracy: 0.6372
Epoch 80, CIFAR-10 Batch 5:	Loss: 0.2192	Validation Accuracy: 0.6426
Epoch 81, CIFAR-10 Batch 1:	Loss: 0.3104	Validation Accuracy: 0.6128
Epoch 81, CIFAR-10 Batch 2:	Loss: 0.2602	Validation Accuracy: 0.6324
Epoch 81, CIFAR-10 Batch 3:	Loss: 0.1870	Validation Accuracy: 0.6450
Epoch 81, CIFAR-10 Batch 4:	Loss: 0.2290	Validation Accuracy: 0.6382
Epoch 81, CIFAR-10 Batch 5:	Loss: 0.1880	Validation Accuracy: 0.6394
Epoch 82, CIFAR-10 Batch 1:	Loss: 0.2928	Validation Accuracy: 0.6054
Epoch 82, CIFAR-10 Batch 2:	Loss: 0.2574	Validation Accuracy: 0.6316
Epoch 82, CIFAR-10 Batch 3:	Loss: 0.2122	Validation Accuracy: 0.6394
Epoch 82, CIFAR-10 Batch 4:	Loss: 0.2171	Validation Accuracy: 0.6400
Epoch 82, CIFAR-10 Batch 5:	Loss: 0.1932	Validation Accuracy: 0.6374
Epoch 83, CIFAR-10 Batch 1:	Loss: 0.2793	Validation Accuracy: 0.6216
Epoch 83, CIFAR-10 Batch 2:	Loss: 0.2431	Validation Accuracy: 0.6240
Epoch 83, CIFAR-10 Batch 3:	Loss: 0.1960	Validation Accuracy: 0.6358
Epoch 83, CIFAR-10 Batch 4:	Loss: 0.2049	Validation Accuracy: 0.6428
Epoch 83, CIFAR-10 Batch 5:	Loss: 0.2085	Validation Accuracy: 0.6498
Epoch 84, CIFAR-10 Batch 1:	Loss: 0.2395	Validation Accuracy: 0.6360
Epoch 84, CIFAR-10 Batch 2:	Loss: 0.2303	Validation Accuracy: 0.6360
Epoch 84, CIFAR-10 Batch 3:	Loss: 0.1980	Validation Accuracy: 0.6394
Epoch 84, CIFAR-10 Batch 4:	Loss: 0.2254	Validation Accuracy: 0.6294
Epoch 84, CIFAR-10 Batch 5:	Loss: 0.1802	Validation Accuracy: 0.6502
Epoch 85, CIFAR-10 Batch 1:	Loss: 0.2700	Validation Accuracy: 0.6276
Epoch 85, CIFAR-10 Batch 2:	Loss: 0.2296	Validation Accuracy: 0.6364
Epoch 85, CIFAR-10 Batch 3:	Loss: 0.1832	Validation Accuracy: 0.6418
Epoch 85, CIFAR-10 Batch 4:	Loss: 0.2096	Validation Accuracy: 0.6388
Epoch 85, CIFAR-10 Batch 5:	Loss: 0.1724	Validation Accuracy: 0.6450
Epoch 86, CIFAR-10 Batch 1:	Loss: 0.2597	Validation Accuracy: 0.6400
Epoch 86, CIFAR-10 Batch 2:	Loss: 0.2264	Validation Accuracy: 0.6378
Epoch 86, CIFAR-10 Batch 3:	Loss: 0.1885	Validation Accuracy: 0.6398
Epoch 86, CIFAR-10 Batch 4:	Loss: 0.2105	Validation Accuracy: 0.6424
Epoch 86, CIFAR-10 Batch 5:	Loss: 0.1747	Validation Accuracy: 0.6490
Epoch 87, CIFAR-10 Batch 1:	Loss: 0.2370	Validation Accuracy: 0.6366
Epoch 87, CIFAR-10 Batch 2:	Loss: 0.2093	Validation Accuracy: 0.6366
Epoch 87, CIFAR-10 Batch 3:	Loss: 0.1741	Validation Accuracy: 0.6472
Epoch 87, CIFAR-10 Batch 4:	Loss: 0.1982	Validation Accuracy: 0.6396
Epoch 87, CIFAR-10 Batch 5:	Loss: 0.1822	Validation Accuracy: 0.6462
Epoch 88, CIFAR-10 Batch 1:	Loss: 0.2474	Validation Accuracy: 0.6386
Epoch 88, CIFAR-10 Batch 2:	Loss: 0.1904	Validation Accuracy: 0.6502
Epoch 88, CIFAR-10 Batch 3:	Loss: 0.1778	Validation Accuracy: 0.6440
Epoch 88, CIFAR-10 Batch 4:	Loss: 0.2058	Validation Accuracy: 0.6468
Epoch 88, CIFAR-10 Batch 5:	Loss: 0.1649	Validation Accuracy: 0.6398
Epoch 89, CIFAR-10 Batch 1:	Loss: 0.2460	Validation Accuracy: 0.6386
Epoch 89, CIFAR-10 Batch 2:	Loss: 0.1920	Validation Accuracy: 0.6458
Epoch 89, CIFAR-10 Batch 3:	Loss: 0.1608	Validation Accuracy: 0.6498
Epoch 89, CIFAR-10 Batch 4:	Loss: 0.1980	Validation Accuracy: 0.6404
Epoch 89, CIFAR-10 Batch 5:	Loss: 0.1560	Validation Accuracy: 0.6434

Epoch 90, CIFAR-10 Batch 1:	Loss: 0.2191	Validation Accuracy: 0.6424
Epoch 90, CIFAR-10 Batch 2:	Loss: 0.1853	Validation Accuracy: 0.6482
Epoch 90, CIFAR-10 Batch 3:	Loss: 0.1737	Validation Accuracy: 0.6468
Epoch 90, CIFAR-10 Batch 4:	Loss: 0.1884	Validation Accuracy: 0.6404
Epoch 90, CIFAR-10 Batch 5:	Loss: 0.1495	Validation Accuracy: 0.6502
Epoch 91, CIFAR-10 Batch 1:	Loss: 0.2226	Validation Accuracy: 0.6442
Epoch 91, CIFAR-10 Batch 2:	Loss: 0.2174	Validation Accuracy: 0.6498
Epoch 91, CIFAR-10 Batch 3:	Loss: 0.1689	Validation Accuracy: 0.6368
Epoch 91, CIFAR-10 Batch 4:	Loss: 0.1676	Validation Accuracy: 0.6434
Epoch 91, CIFAR-10 Batch 5:	Loss: 0.1546	Validation Accuracy: 0.6392
Epoch 92, CIFAR-10 Batch 1:	Loss: 0.2123	Validation Accuracy: 0.6364
Epoch 92, CIFAR-10 Batch 2:	Loss: 0.1678	Validation Accuracy: 0.6388
Epoch 92, CIFAR-10 Batch 3:	Loss: 0.1565	Validation Accuracy: 0.6364
Epoch 92, CIFAR-10 Batch 4:	Loss: 0.1873	Validation Accuracy: 0.6390
Epoch 92, CIFAR-10 Batch 5:	Loss: 0.1732	Validation Accuracy: 0.6428
Epoch 93, CIFAR-10 Batch 1:	Loss: 0.2346	Validation Accuracy: 0.6336
Epoch 93, CIFAR-10 Batch 2:	Loss: 0.1776	Validation Accuracy: 0.6410
Epoch 93, CIFAR-10 Batch 3:	Loss: 0.1882	Validation Accuracy: 0.6382
Epoch 93, CIFAR-10 Batch 4:	Loss: 0.1949	Validation Accuracy: 0.6378
Epoch 93, CIFAR-10 Batch 5:	Loss: 0.1514	Validation Accuracy: 0.6416
Epoch 94, CIFAR-10 Batch 1:	Loss: 0.2407	Validation Accuracy: 0.6344
Epoch 94, CIFAR-10 Batch 2:	Loss: 0.1696	Validation Accuracy: 0.6326
Epoch 94, CIFAR-10 Batch 3:	Loss: 0.1772	Validation Accuracy: 0.6340
Epoch 94, CIFAR-10 Batch 4:	Loss: 0.1968	Validation Accuracy: 0.6410
Epoch 94, CIFAR-10 Batch 5:	Loss: 0.1820	Validation Accuracy: 0.6410
Epoch 95, CIFAR-10 Batch 1:	Loss: 0.2244	Validation Accuracy: 0.6336
Epoch 95, CIFAR-10 Batch 2:	Loss: 0.1742	Validation Accuracy: 0.6320
Epoch 95, CIFAR-10 Batch 3:	Loss: 0.1717	Validation Accuracy: 0.6316
Epoch 95, CIFAR-10 Batch 4:	Loss: 0.1783	Validation Accuracy: 0.6318
Epoch 95, CIFAR-10 Batch 5:	Loss: 0.1520	Validation Accuracy: 0.6384
Epoch 96, CIFAR-10 Batch 1:	Loss: 0.2517	Validation Accuracy: 0.6302
Epoch 96, CIFAR-10 Batch 2:	Loss: 0.1708	Validation Accuracy: 0.6336
Epoch 96, CIFAR-10 Batch 3:	Loss: 0.1789	Validation Accuracy: 0.6248
Epoch 96, CIFAR-10 Batch 4:	Loss: 0.1887	Validation Accuracy: 0.6356
Epoch 96, CIFAR-10 Batch 5:	Loss: 0.1775	Validation Accuracy: 0.6246
Epoch 97, CIFAR-10 Batch 1:	Loss: 0.2304	Validation Accuracy: 0.6290
Epoch 97, CIFAR-10 Batch 2:	Loss: 0.1752	Validation Accuracy: 0.6338
Epoch 97, CIFAR-10 Batch 3:	Loss: 0.1847	Validation Accuracy: 0.6234
Epoch 97, CIFAR-10 Batch 4:	Loss: 0.2100	Validation Accuracy: 0.6440
Epoch 97, CIFAR-10 Batch 5:	Loss: 0.1756	Validation Accuracy: 0.6250
Epoch 98, CIFAR-10 Batch 1:	Loss: 0.2312	Validation Accuracy: 0.6346
Epoch 98, CIFAR-10 Batch 2:	Loss: 0.1651	Validation Accuracy: 0.6380
Epoch 98, CIFAR-10 Batch 3:	Loss: 0.1646	Validation Accuracy: 0.6332
Epoch 98, CIFAR-10 Batch 4:	Loss: 0.2051	Validation Accuracy: 0.6398
Epoch 98, CIFAR-10 Batch 5:	Loss: 0.1639	Validation Accuracy: 0.6170
Epoch 99, CIFAR-10 Batch 1:	Loss: 0.2093	Validation Accuracy: 0.6452
Epoch 99, CIFAR-10 Batch 2:	Loss: 0.1753	Validation Accuracy: 0.6388
Epoch 99, CIFAR-10 Batch 3:	Loss: 0.1700	Validation Accuracy: 0.6346


```
Epoch 99, CIFAR-10 Batch 4: Loss: 0.1869 Validation Accuracy: 0.6384
Epoch 99, CIFAR-10 Batch 5: Loss: 0.1447 Validation Accuracy: 0.6144
Epoch 100, CIFAR-10 Batch 1: Loss: 0.2100 Validation Accuracy: 0.6448
Epoch 100, CIFAR-10 Batch 2: Loss: 0.1642 Validation Accuracy: 0.6296
Epoch 100, CIFAR-10 Batch 3: Loss: 0.1953 Validation Accuracy: 0.6218
Epoch 100, CIFAR-10 Batch 4: Loss: 0.1992 Validation Accuracy: 0.6394
Epoch 100, CIFAR-10 Batch 5: Loss: 0.1527 Validation Accuracy: 0.6220
```

3 Checkpoint

The model has been saved to disk. ## Test Model Test your model against the test dataset. This will be your final accuracy. You should have an accuracy greater than 50%. If you don't, keep tweaking the model architecture and parameters.

```
In [193]: """
          DON'T MODIFY ANYTHING IN THIS CELL
          """

          %matplotlib inline
          %config InlineBackend.figure_format = 'retina'

          import tensorflow as tf
          import pickle
          import helper
          import random

          # Set batch size if not already set
          try:
              if batch_size:
                  pass
          except NameError:
              batch_size = 64

          save_model_path = './image_classification'
          n_samples = 4
          top_n_predictions = 3

          def test_model():
              """
              Test the saved model against the test dataset
              """

              test_features, test_labels = pickle.load(open('preprocess_training.p', mode='rb'))
              loaded_graph = tf.Graph()

              with tf.Session(graph=loaded_graph) as sess:
                  # Load model
```

```

loader = tf.train.import_meta_graph(save_model_path + '.meta')
loader.restore(sess, save_model_path)

# Get Tensors from loaded model
loaded_x = loaded_graph.get_tensor_by_name('x:0')
loaded_y = loaded_graph.get_tensor_by_name('y:0')
loaded_keep_prob = loaded_graph.get_tensor_by_name('keep_prob:0')
loaded_logits = loaded_graph.get_tensor_by_name('logits:0')
loaded_acc = loaded_graph.get_tensor_by_name('accuracy:0')

# Get accuracy in batches for memory limitations
test_batch_acc_total = 0
test_batch_count = 0

for train_feature_batch, train_label_batch in helper.batch_features_labels(test_data_loader):
    test_batch_acc_total += sess.run(
        loaded_acc,
        feed_dict={loaded_x: train_feature_batch, loaded_y: train_label_batch,
                    loaded_keep_prob: 1.0})
    test_batch_count += 1

print('Testing Accuracy: {}'.format(test_batch_acc_total/test_batch_count))

# Print Random Samples
random_test_features, random_test_labels = tuple(zip(*random.sample(list(zip(test_data_loader.get_test_features(),
test_data_loader.get_test_labels()), 10))))
random_test_predictions = sess.run(
    tf.nn.top_k(tf.nn.softmax(loaded_logits), top_n_predictions),
    feed_dict={loaded_x: random_test_features, loaded_y: random_test_labels,
                loaded_keep_prob: 1.0})
helper.display_image_predictions(random_test_features, random_test_labels, random_test_predictions)


```

```
test_model()
```


Testing Accuracy: 0.61728515625

Softmax Predictions


truck





horse



bird



deer



3.1 Why 50-80% Accuracy?

You might be wondering why you can't get an accuracy any higher. First things first, 50% isn't bad for a simple CNN. Pure guessing would get you 10% accuracy. However, you might notice people are getting scores [well above 80%](#). That's because we haven't taught you all there is to know about neural networks. We still need to cover a few more techniques. ## Submitting This Project When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd_image_classification.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "helper.py" and "problem_unittests.py" files in your submission.